

FRAUD DETECTION IN A GOVERNMENT AGENCY (MLOPS)

INTRODUCTION

- **CONTEXT:** Increasing number of fraud suspicions in government agency program
- **PROBLEM:** Thousands of applications, manual fraud detection not possible
- **GOAL:** Automated fraud detection system, adaptable to changing data
- **APPROACH:** Apply MLOps principles like iterative model training and retraining, monitoring, experiment logging and versioning (Luu et al., 2024, pp. 2-4)
- **DATASET:** Credit Card Fraud Detection (Credit Card Fraud Detection, 2025)

OUTLINE

Requirement Analysis

System Architecture

Fraud Detection Model

MLflow Tracking & Logging with S3 Backups

Retraining Scenarios

REST-API

Simulation of One Year

Challenges and Limitations

Conclusion

REQUIREMENT ANALYSIS




Automated fraud
detection



Easily adaptable
and retrainable



Monitored for
reliability and
performance



RESTful API
serves
integration



Secure Acces
Control



Cloud-ready /
scalable

SYSTEM ARCHITECTURE

Data Flow:

- Online applications → weekly CSVs stored in S3 → Drift Watchdog → Retraining → Random Forest → Flask API

Docker Container:

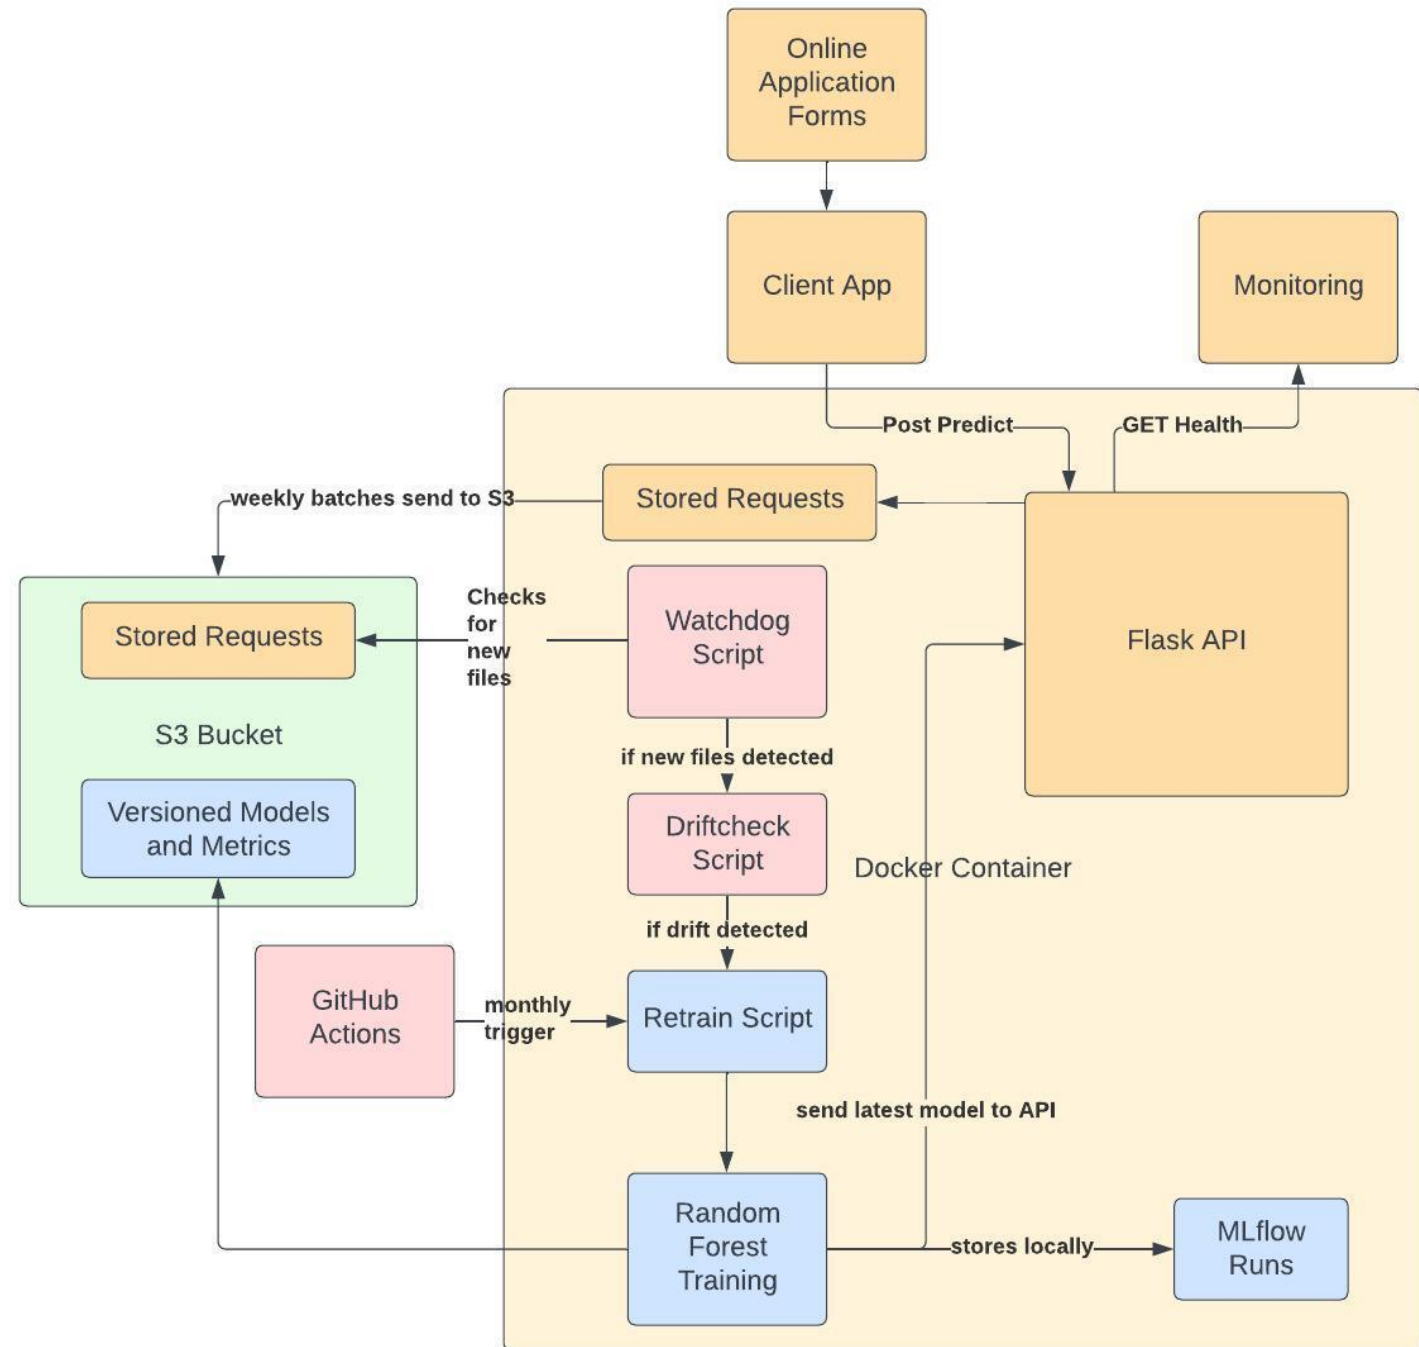
- hosts scripts & API
- ensures consistent environments (McKendrick & Gallagher, 2018, pp. 4–8)

MLOps:

- GitHub Actions automates retraining

Security:

- API token for access control



FRAUD DETECTION MODEL: RANDOM FOREST

- “Slim” tree for faster training & retraining
- Robust, handles imbalanced data, reduces false negatives (Mihali & Niță, 2024, p.111)
- Tuned hyperparameters: 25 trees, max depth 8, min samples split 5, min samples leaf 3
- Trained on the last 12 weeks of data from S3
- Predicts fraud probability for new applications

MODEL VERSIONING AND LOGGING

- **MLflow:**
 - tracks experiments inside Docker Container
 - Logs training metrics, model artifacts (Luu et al., 2024, pp. 136-150)
- **S3-Backup:**
 - Models and metrics versioned and stored in S3
- → This Setup enables performance monitoring

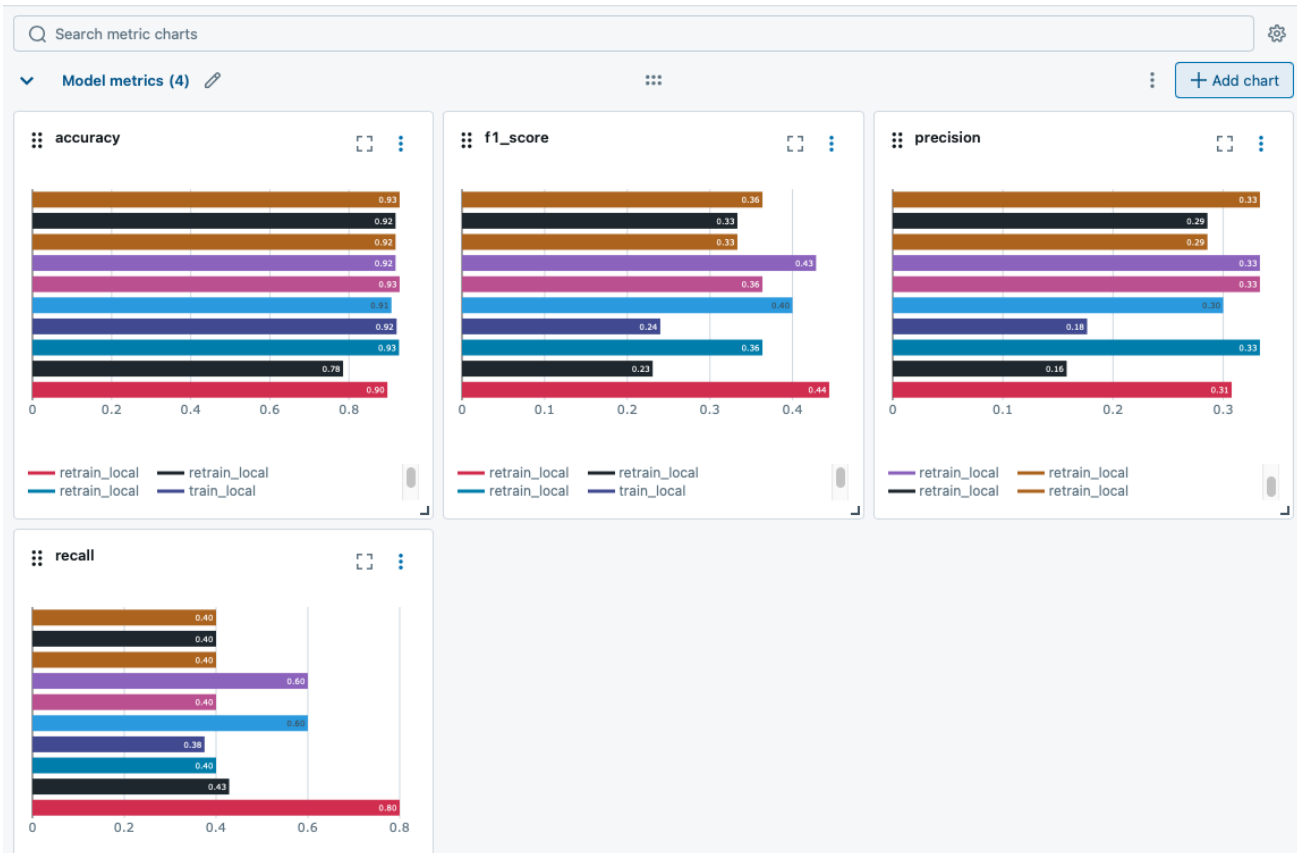


Image Source: Author

SCHEDULED MONTHLY RETRAINING

- Retraining triggered automatically every month via **GitHub Actions**
- Uses the last four weeks of data from S3
- New model logged in MLflow and backed up in S3 and deployed to Flask API
- Ensures system adapts to evolving patterns over time

← Monthly Model Retraining

✓ Monthly Model Retraining #5

🏠 Summary

Jobs

✓ retrain

Run details

🕒 Usage

📄 Workflow file

retrain

succeeded now in 16m 31s

- > ✓ Set up job
- > ✓ Checkout repository
- > ✓ Set up Docker
- > ✓ Build Docker image
- > ✓ Run retraining container for 15 minutes
- > ✓ Verify MLflow logs
- > ✓ Post Set up Docker
- > ✓ Post Checkout repository
- > ✓ Complete job

DRIFT BASED RETRAINING

- All API requests are stored locally and uploaded weekly to S3.
- **Drift Watchdog script** detects new files
- For each new file:
 - Compare feature distributions using two-sample-t-test (Kutner et al. 2013, pp 278-288)
 - Bonferroni correction applied: α divided by number of features to control overall false positives (Kutner et al. 2013, pp 155-156)
 - Corrected p-value $< 0.05 \rightarrow$ feature marked as drifted
- **Drift detected** \rightarrow Trigger retraining in Docker
- **MLflow** logs new model and metrics
- **S3 backup** stores versioned model, metrics, input example
- **API** updated with latest model

SIMULATION OF ONE YEAR

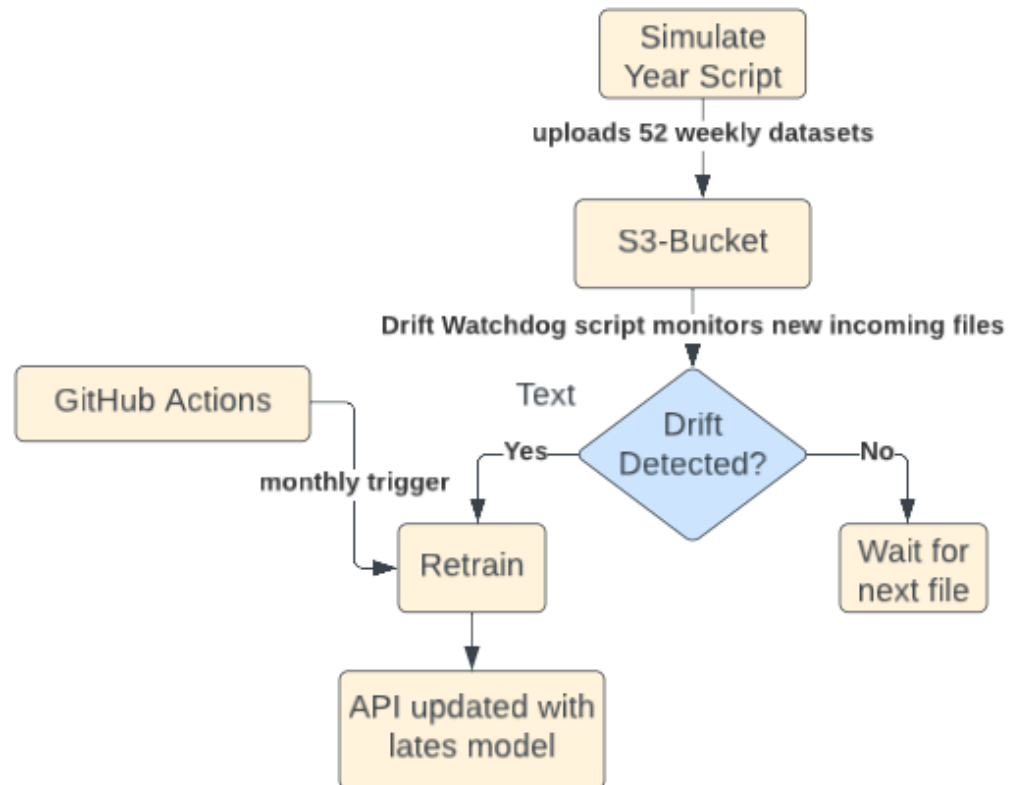


Image Source: Author

- Simulation script continuously uploads new files to S3
- Drift Watchdog script runs in Docker, monitoring new files in real-time
- Validates that MLflow logging, S3 backups, and API updates happen as expected


REST-API

Available Actions:

- **POST** /predict → submit new application data, get fraud probability
- **GET** /health → check API and system status

Process Flow:

- Client sends request to the Flask API
- API accesses the latest Random Forest model
- Model returns predictions as JSON



CHALLENGES AND LIMITATIONS

Challenges:

- S3 connection: credentials & paths
- Deploying on EC2 was challenging

Limitations & Next Steps:

- **Scaling & Deployment:**

- Single container limits scalability and security
→ consider Docker Swarm for better isolation and encrypted communication (Ganne, 2022, p.4)

- **Compute & Performance:**

- t3.micro EC2 is slow for training/retraining → upgrade instance

- **Model & Monitoring:**

- Simple drift detection, small Random Forest
→ implement advanced drift detection, consider larger ensemble models

CONCLUSION

- System successfully detects fraud automatically using Random Forest
- Retraining triggered monthly or on data drift
- Models logged in MLflow and versioned backups stored in S3
- REST API provides seamless integration with client applications
- **Next Steps:** Deploy on EC2, consider Docker Swarm, enhance monitoring

BIBLIOGRAPHY

Credit Card Fraud Detection. (2025). [Dataset]. <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>

Ganne, A. (2022). CLOUD DATA SECURITY METHODS: KUBERNETES VS DOCKER SWARM. *International Research Journal of Modernization in Engineering Technology and Science*. <https://doi.org/10.56726/IRJMETS32176>

Kutner, M. H., Nachtsheim, C. J., Neter, J., & Li, W. (2013). *Applied Linear Statistical Models* (5th ed.).

Luu, H., Pumperla, M., & Zhang, Z. (2024). *MLOps with Ray: Best practices and strategies for adopting machine learning operations*. Apress L. P.

McKendrick, R., & Gallagher, S. (2018). *Mastering docker: Unlock new opportunities using Docker's most advanced features, third edition* (3rd ed). Packt Publishing.

Mihali, S.-I., & Niță, Ștefania-L. (2024). Credit Card Fraud Detection based on Random Forest Model. *2024 International Conference on Development and Application Systems (DAS)*, 111–114. <https://doi.org/10.1109/DAS61944.2024.10541240>

LIST OF FIGURES

- Figure 1: System Architecture, p. 5 (source: Author)
- Figure 2: Model Logging, p. 7 (source: Author)
- Figure 3: GitHub Actions Monthly Trigger, p. 15 (source: Author)
- Figure 4: Simulation of One Year, p. 9 (source: Author)